

House Prices : Data cleaning, visualization and modeling

In [1]:

```
# Import packages
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn import preprocessing
from sklearn.model_selection import train_test_split
from lightgbm import LGBMRegressor
from xgboost import XGBRegressor
import sklearn.metrics as metrics
import math
```

Importing **train** and **test** datasets

In [2]:

```
sample_submission = pd.read_csv("../input/house-prices-advanced-regression-techniques/sample_submission.csv")
test = pd.read_csv("../input/house-prices-advanced-regression-techniques/test.csv")
train = pd.read_csv("../input/house-prices-advanced-regression-techniques/train.csv")
#Creating a copy of the train and test datasets
c_test = test.copy()
c_train = train.copy()
```

- Getting information about train dataset

In [3]:

```
c_train.head()
```

Out[3]:

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	LandContour	U1
0	1	60	RL	65.0	8450	Pave	NaN	Reg	Lvl	/
1	2	20	RL	80.0	9600	Pave	NaN	Reg	Lvl	/
2	3	60	RL	68.0	11250	Pave	NaN	IR1	Lvl	/
3	4	70	RL	60.0	9550	Pave	NaN	IR1	Lvl	/
4	5	60	RL	84.0	14260	Pave	NaN	IR1	Lvl	/

5 rows × 81 columns

- Getting information about test dataset

In [4]:

```
c_test.head()
```

Out[4]:

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	LandContour
0	1461	20	RH	80.0	11622	Pave	NaN	Reg	Lvl
1	1462	20	RL	81.0	14267	Pave	NaN	IR1	Lvl
2	1463	60	RL	74.0	13830	Pave	NaN	IR1	Lvl
3	1464	60	RL	78.0	9978	Pave	NaN	IR1	Lvl
4	1465	120	RL	43.0	5005	Pave	NaN	IR1	HLS

5 rows × 80 columns

1. We have 81 columns. 2. Our target variable is SalePrice. 3. Id is just an index that we can drop but we will need it in the final submission.

- 1. We have many missing values

* * * * we have 79 features in our dataset.

- Concat Train and Test datasets

In [5]:

```
c_train['train'] = 1
c_test['train'] = 0
df = pd.concat([c_train, c_test], axis=0, sort=False)
```

Data preprocessing

- Calculating the percentage of missing values of each feature

In [6]:

```
#Percentage of NAN Values
NAN = [(c, df[c].isna().mean()*100) for c in df]
NAN = pd.DataFrame(NAN, columns=["column_name", "percentage"])
```

- Features with more than 50% of missing values.

In [7]:

```
NAN = NAN[NAN.percentage > 50]
NAN.sort_values("percentage", ascending=False)
```

Out[7]:

	column_name	percentage
72	PoolQC	99.657417
74	MiscFeature	96.402878
6	Alley	93.216855
73	Fence	80.438506

- We can drop PoolQC, MiscFeature, Alley and Fence features because they have more than 80% of missing values.

In [8]:

```
#Drop PoolQC, MiscFeature, Alley and Fence features
df = df.drop(['Alley', 'PoolQC', 'Fence', 'MiscFeature'], axis=1)
```

- Now we will select numerical and categorical features

In [9]:

```
object_columns_df = df.select_dtypes(include=['object'])
numerical_columns_df = df.select_dtypes(exclude=['object'])
```

- **Categorical Features :**

In [10]:

```
object_columns_df.dtypes
```

Out[10]:

MSZoning	object
Street	object
LotShape	object
LandContour	object
Utilities	object
LotConfig	object
LandSlope	object
Neighborhood	object
Condition1	object
Condition2	object
BldgType	object
HouseStyle	object
RoofStyle	object
RoofMatl	object
Exterior1st	object
Exterior2nd	object
MasVnrType	object
ExterQual	object
ExterCond	object
Foundation	object
BsmtQual	object
BsmtCond	object
BsmtExposure	object
BsmtFinType1	object
BsmtFinType2	object
Heating	object
HeatingQC	object
CentralAir	object
Electrical	object
KitchenQual	object
Functional	object
FireplaceQu	object
GarageType	object
GarageFinish	object
GarageQual	object
GarageCond	object
PavedDrive	object
SaleType	object
SaleCondition	object
	dtype: object

- **Numerical Features :**

In [11]:

```
numerical_columns_df.dtypes
```

Out[11]:

```
Id           int64
MSSubClass   int64
LotFrontage  float64
LotArea      int64
OverallQual  int64
OverallCond  int64
YearBuilt    int64
YearRemodAdd int64
MasVnrArea   float64
BsmtFinSF1  float64
BsmtFinSF2  float64
BsmtUnfSF   float64
TotalBsmtSF float64
1stFlrSF     int64
2ndFlrSF     int64
LowQualFinSF int64
GrLivArea    int64
BsmtFullBath float64
BsmtHalfBath float64
FullBath     int64
HalfBath     int64
BedroomAbvGr int64
KitchenAbvGr int64
TotRmsAbvGrd int64
Fireplaces   int64
GarageYrBlt  float64
GarageCars   float64
GarageArea   float64
WoodDeckSF   int64
OpenPorchSF  int64
EnclosedPorch int64
3SsnPorch    int64
ScreenPorch  int64
PoolArea     int64
MiscVal      int64
MoSold       int64
YrSold       int64
SalePrice    float64
train        int64
dtype: object
```

- Dealing with **categorical** feature

In [12]:

```
#Number of null values in each feature
null_counts = object_columns_df.isnull().sum()
print("Number of null values in each column:\n{}".format(null_counts))
```

Number of null values in each column:

MSZoning	4
Street	0
LotShape	0
LandContour	0
Utilities	2
LotConfig	0
LandSlope	0
Neighborhood	0
Condition1	0
Condition2	0
BldgType	0
HouseStyle	0
RoofStyle	0
RoofMatl	0
Exterior1st	1
Exterior2nd	1
MasVnrType	24
ExterQual	0
ExterCond	0
Foundation	0
BsmtQual	81
BsmtCond	82
BsmtExposure	82
BsmtFinType1	79
BsmtFinType2	80
Heating	0
HeatingQC	0
CentralAir	0
Electrical	1
KitchenQual	1
Functional	2
FireplaceQu	1420
GarageType	157
GarageFinish	159
GarageQual	159
GarageCond	159
PavedDrive	0
SaleType	1
SaleCondition	0

dtype: int64

- We will fill -- **BsmtQual**, **BsmtCond**, **BsmtExposure**, **BsmtFinType1**, **BsmtFinType2**, **GarageType**, **GarageFinish**, **GarageQual**, **FireplaceQu**, **GarageCond** -- with "None" (Take a look in the data description).
- We will fill the rest of features with th most frequent value (using its own most frequent value).

In [13]:

```
columns_None = ['BsmtQual', 'BsmtCond', 'BsmtExposure', 'BsmtFinType1', 'BsmtFinType2', 'GarageType', 'GarageFinish', 'GarageCond', 'PorchType', 'PorchStyle', 'Fence', 'MiscFeature']
object_columns_df[columns_None] = object_columns_df[columns_None].fillna('None')
```

/opt/conda/lib/python3.6/site-packages/pandas/core/frame.py:3509: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using `.loc[row_indexer, col_indexer] = value` instead

See the caveats in the documentation: http://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (http://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

`self[k1] = value[k2]`

In [14]:

```
columns_with_lowNA = ['MSZoning', 'Utilities', 'Exterior1st', 'Exterior2nd', 'MasVnrType', 'Electrical',  
#fill missing values for each column (using its own most frequent value)  
object_columns_df[columns_with_lowNA] = object_columns_df[columns_with_lowNA].fillna(object_columns_
```

- Now we have a clean categorical features
 - In the next step we will deal with the **numerical** features black

In [15]:

```
#Number of null values in each feature
null_counts = numerical_columns_df.isnull().sum()
print("Number of null values in each column:\n{}".format(null_counts))
```

Number of null values in each column:

Id	0
MSSubClass	0
LotFrontage	486
LotArea	0
OverallQual	0
OverallCond	0
YearBuilt	0
YearRemodAdd	0
MasVnrArea	23
BsmtFinSF1	1
BsmtFinSF2	1
BsmtUnfSF	1
TotalBsmtSF	1
1stFlrSF	0
2ndFlrSF	0
LowQualFinSF	0
GrLivArea	0
BsmtFullBath	2
BsmtHalfBath	2
FullBath	0
HalfBath	0
BedroomAbvGr	0
KitchenAbvGr	0
TotRmsAbvGrd	0
Fireplaces	0
GarageYrBlt	159
GarageCars	1
GarageArea	1
WoodDeckSF	0
OpenPorchSF	0
EnclosedPorch	0
3SsnPorch	0
ScreenPorch	0
PoolArea	0
MiscVal	0
MoSold	0
YrSold	0
SalePrice	1459
train	0

dtype: int64

1. Fill GarageYrBlt and LotFrontage
2. Fill the rest of columns with 0

In [16]:

```
print((numerical_columns_df['YrSold']-numerical_columns_df['YearBuilt']).median())
print(numerical_columns_df["LotFrontage"].median())
```

35.0
68.0

So we will fill the year with 1979 and the Lot frontage with 68

In [17]:

```
numerical_columns_df['GarageYrBlt'] = numerical_columns_df['GarageYrBlt'].fillna(numerical_columns_df['GarageYrBlt'].median())
numerical_columns_df['LotFrontage'] = numerical_columns_df['LotFrontage'].fillna(68)
```

/opt/conda/lib/python3.6/site-packages/ipykernel_launcher.py:1: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: http://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

"""Entry point for launching an IPython kernel.

/opt/conda/lib/python3.6/site-packages/ipykernel_launcher.py:2: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: http://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

Fill the rest of columns with 0

In [18]:

```
numerical_columns_df = numerical_columns_df.fillna(0)
```


- We finally end up with a clean dataset
- After making some plots we found that we have some columns with low variance so we decide to delete them

In [19]:

```
object_columns_df['Utilities'].value_counts().plot(kind='bar', figsize=[10, 3])
object_columns_df['Utilities'].value_counts()
```

Out[19]:

```
AllPub    2918
NoSeWa      1
Name: Utilities, dtype: int64
```




In [20]:

```
object_columns_df['Street'].value_counts().plot(kind='bar', figsize=[10, 3])
object_columns_df['Street'].value_counts()
```

Out[20]:

```
Pave    2907
Grvl     12
Name: Street, dtype: int64
```




In [21]:

```
object_columns_df['Condition2'].value_counts().plot(kind='bar', figsize=[10, 3])  
object_columns_df['Condition2'].value_counts()
```

Out[21]:

```
Norm      2889  
Feedr     13  
Artery     5  
PosA       4  
PosN       4  
RRNn       2  
RRAe       1  
RRAn       1  
Name: Condition2, dtype: int64
```




In [22]:

```
object_columns_df['RoofMatl'].value_counts().plot(kind='bar', figsize=[10, 3])
object_columns_df['RoofMatl'].value_counts()
```

Out[22]:

CompShg	2876
Tar&Grv	23
WdShake	9
WdShngl	7
ClyTile	1
Metal	1
Membran	1
Roll	1

Name: RoofMatl, dtype: int64




In [23]:

```
object_columns_df['Heating'].value_counts().plot(kind='bar', figsize=[10, 3])
object_columns_df['Heating'].value_counts() #=====> Drop feature one Type
```

Out[23]:

GasA	2874
GasW	27
Grav	9
Wall	6
OthW	2
Floor	1

Name: Heating, dtype: int64



In [24]:

```
object_columns_df = object_columns_df.drop(['Heating', 'RoofMatl', 'Condition2', 'Street', 'Utilities'],
```

- Now we will create some new features

In [25]:

```
numerical_columns_df['Age_House'] = (numerical_columns_df['YrSold'] - numerical_columns_df['YearBuilt'])  
numerical_columns_df['Age_House'].describe()
```

Out[25]:

```
count    2919.000000  
mean     36.479959  
std      30.336182  
min     -1.000000  
25%      7.000000  
50%     35.000000  
75%     54.500000  
max     136.000000  
Name: Age_House, dtype: float64
```


In [26]:

```
Negatif = numerical_columns_df[numerical_columns_df['Age_House'] < 0]  
Negatif
```

Out[26]:

	Id	MSSubClass	LotFrontage	LotArea	OverallQual	OverallCond	YearBuilt	YearRemo
1089	2550	20	128.0	39290	10	5	2008	

1 rows × 40 columns



- Like we see here tha the minimun is -1 ???
- It is strange to find that the house was sold in 2007 before the YearRemodAdd 2009.

So we decide to change the year of sold to 2009

In [27]:

```
numerical_columns_df.loc[numerical_columns_df['YrSold'] < numerical_columns_df['YearBuilt'], 'YrSold'] = numerical_columns_df['Age_House'] = (numerical_columns_df['YrSold'] - numerical_columns_df['YearBuilt'])
numerical_columns_df['Age_House'].describe()
```

Out[27]:

count	2919.000000
mean	36.480644
std	30.335358
min	0.000000
25%	7.000000
50%	35.000000
75%	54.500000
max	136.000000
Name:	Age_House, dtype: float64

- * **TotalBsmtBath** : Sum of : BsmtFullBath and 1/2 BsmtHalfBath

- **TotalBath** : Sum of : FullBath and 1/2 HalfBath
- **TotalSA** : Sum of : 1stFlrSF and 2ndFlrSF and basement area

In [28]:

```
numerical_columns_df['TotalBsmtBath'] = numerical_columns_df['BsmtFullBath'] + numerical_columns_df['BsmtHalfBath']
numerical_columns_df['TotalBath'] = numerical_columns_df['FullBath'] + numerical_columns_df['HalfBath']
numerical_columns_df['TotalSA'] = numerical_columns_df['TotalBsmtSF'] + numerical_columns_df['1stFlrSF'] + numerical_columns_df['2ndFlrSF']
```


In [29]:

```
numerical_columns_df.head()
```

Out[29]:

	Id	MSSubClass	LotFrontage	LotArea	OverallQual	OverallCond	YearBuilt	YearRemodAdd
0	1	60	65.0	8450	7	5	2003	2003
1	2	20	80.0	9600	6	8	1976	1976
2	3	60	68.0	11250	7	5	2001	2002
3	4	70	60.0	9550	7	5	1915	1970
4	5	60	84.0	14260	8	5	2000	2000

5 rows × 43 columns



- Now the next step is to encode categorical features
- **Ordinal categories features** - Mapping from 0 to N

In [30]:

```

bin_map = {'TA': 2, 'Gd': 3, 'Fa': 1, 'Ex': 4, 'Po': 1, 'None': 0, 'Y': 1, 'N': 0, 'Reg': 3, 'IR1': 2, 'IR2': 1, 'IR3': 0,
           'No': 2, "Mn": 2, "Av": 3, "Gd": 4, "Unf": 1, "LwQ": 2, "Rec": 3, "BLQ": 4, "ALQ": 5
         }

object_columns_df['ExterQual'] = object_columns_df['ExterQual'].map(bin_map)
object_columns_df['ExterCond'] = object_columns_df['ExterCond'].map(bin_map)
object_columns_df['BsmtCond'] = object_columns_df['BsmtCond'].map(bin_map)
object_columns_df['BsmtQual'] = object_columns_df['BsmtQual'].map(bin_map)
object_columns_df['HeatingQC'] = object_columns_df['HeatingQC'].map(bin_map)
object_columns_df['KitchenQual'] = object_columns_df['KitchenQual'].map(bin_map)
object_columns_df['FireplaceQu'] = object_columns_df['FireplaceQu'].map(bin_map)
object_columns_df['GarageQual'] = object_columns_df['GarageQual'].map(bin_map)
object_columns_df['GarageCond'] = object_columns_df['GarageCond'].map(bin_map)
object_columns_df['CentralAir'] = object_columns_df['CentralAir'].map(bin_map)
object_columns_df['LotShape'] = object_columns_df['LotShape'].map(bin_map)
object_columns_df['BsmtExposure'] = object_columns_df['BsmtExposure'].map(bin_map)
object_columns_df['BsmtFinType1'] = object_columns_df['BsmtFinType1'].map(bin_map)
object_columns_df['BsmtFinType2'] = object_columns_df['BsmtFinType2'].map(bin_map)

PavedDrive = {"N": 0, "P": 1, "Y": 2}
object_columns_df['PavedDrive'] = object_columns_df['PavedDrive'].map(PavedDrive)

```

- Will we use One hot encoder to encode the rest of categorical features

In [31]:

```

#Select categorical features
rest_object_columns = object_columns_df.select_dtypes(include=['object'])
#Using One hot encoder
object_columns_df = pd.get_dummies(object_columns_df, columns=rest_object_columns.columns)

```

In [32]:

```
object_columns_df.head()
```

Out[32]:

	LotShape	ExterQual	ExterCond	BsmtQual	BsmtCond	BsmtExposure	BsmtFinType1	Bsmt
0	3	4	2	4	2	2		6
1	3	2	2	4	2	4		5
2	2	4	2	4	2	2		6
3	2	2	2	2	4	2		5
4	2	4	2	4	2	3		6

5 rows × 164 columns

- Concat Categorical (after encoding) and numerical features

In [33]:

```
df_final = pd.concat([object_columns_df, numerical_columns_df], axis=1, sort=False)
df_final.head()
```

Out[33]:

	LotShape	ExterQual	ExterCond	BsmtQual	BsmtCond	BsmtExposure	BsmtFinType1	Bsmt
0	3	4	2	4	2	2		6
1	3	2	2	4	2	4		5
2	2	4	2	4	2	2		6
3	2	2	2	2	4	2		5
4	2	4	2	4	2	3		6

5 rows × 207 columns

In [34]:

```
df_final = df_final.drop(['Id'], axis=1)

df_train = df_final[df_final['train'] == 1]
df_train = df_train.drop(['train'], axis=1)

df_test = df_final[df_final['train'] == 0]
df_test = df_test.drop(['SalePrice'], axis=1)
df_test = df_test.drop(['train'], axis=1)
```

- Separate Train and Targets

In [35]:

```
target = df_train['SalePrice']
df_train = df_train.drop(['SalePrice'], axis=1)
```

Modeling

In [36]:

```
x_train, x_test, y_train, y_test = train_test_split(df_train, target, test_size=0.33, random_state=0)
```

In [37]:

```
xgb = XGBRegressor( booster='gbtree', colsample_bylevel=1,
                     colsample_bynode=1, colsample_bytree=0.6, gamma=0,
                     importance_type='gain', learning_rate=0.01, max_delta_step=0,
                     max_depth=4, min_child_weight=1.5, n_estimators=2400,
                     n_jobs=1, nthread=None, objective='reg:linear',
                     reg_alpha=0.6, reg_lambda=0.6, scale_pos_weight=1,
                     silent=None, subsample=0.8, verbosity=1)

lgbm = LGBMRegressor(objective='regression',
                      num_leaves=4,
                      learning_rate=0.01,
                      n_estimators=12000,
                      max_bin=200,
                      bagging_fraction=0.75,
                      bagging_freq=5,
                      bagging_seed=7,
                      feature_fraction=0.4,
                      )
```

In [38]:

```
#Fitting
xgb.fit(x_train, y_train)
lgbm.fit(x_train, y_train, eval_metric='rmse')
```

/opt/conda/lib/python3.6/site-packages/xgboost/core.py:587: FutureWarning: Series.base is deprecated and will be removed in a future version
 if getattr(data, 'base', None) is not None and \

[22:45:26] WARNING: /workspace/src/objective/regression_obj.cu:152: reg:linear is now deprecated in favor of reg:squarederror.

Out[38]:

```
LGBMRegressor(bagging_fraction=0.75, bagging_freq=5, bagging_seed=7,
               boosting_type='gbdt', class_weight=None, colsample_bytree=1.0,
               feature_fraction=0.4, importance_type='split', learning_rate=0.01,
               max_bin=200, max_depth=-1, min_child_samples=20,
               min_child_weight=0.001, min_split_gain=0.0, n_estimators=12000,
               n_jobs=-1, num_leaves=4, objective='regression',
               random_state=None, reg_alpha=0.0, reg_lambda=0.0, silent=True,
               subsample=1.0, subsample_for_bin=200000, subsample_freq=0)
```

In [39]:

```
predict1 = xgb.predict(x_test)
predict = lgbm.predict(x_test)
```

In [40]:

```
print('Root Mean Square Error test = ' + str(math.sqrt(metrics.mean_squared_error(y_test, predict1)))
print('Root Mean Square Error test = ' + str(math.sqrt(metrics.mean_squared_error(y_test, predict))))
```

Root Mean Square Error test = 26089.30803984919
Root Mean Square Error test = 26726.198278903514

- Fitting With all the dataset

In [41]:

```
xgb.fit(df_train, target)
lgbm.fit(df_train, target, eval_metric='rmse')
```

[22:45:47] WARNING: /workspace/src/objective/regression_obj.cu:152: reg:linear is now deprecated in favor of reg:squarederror.

```
/opt/conda/lib/python3.6/site-packages/xgboost/core.py:587: FutureWarning: Series.base is deprecated and will be removed in a future version
    if getattr(data, 'base', None) is not None and \
/opt/conda/lib/python3.6/site-packages/xgboost/core.py:588: FutureWarning: Series.base is deprecated and will be removed in a future version
    data.base is not None and isinstance(data, np.ndarray) \
```

Out[41]:

```
LGBMRegressor(bagging_fraction=0.75, bagging_freq=5, bagging_seed=7,
               boosting_type='gbdt', class_weight=None, colsample_bytree=1.0,
               feature_fraction=0.4, importance_type='split', learning_rate=0.01,
               max_bin=200, max_depth=-1, min_child_samples=20,
               min_child_weight=0.001, min_split_gain=0.0, n_estimators=12000,
               n_jobs=-1, num_leaves=4, objective='regression',
               random_state=None, reg_alpha=0.0, reg_lambda=0.0, silent=True,
               subsample=1.0, subsample_for_bin=200000, subsample_freq=0)
```

In [42]:

```
predict4 = lgbm.predict(df_test)
predict3 = xgb.predict(df_test)
predict_y = (predict3*0.45 + predict4 * 0.55)
```

In [43]:

```
submission = pd.DataFrame({
    "Id": test["Id"],
    "SalePrice": predict_y
})
submission.to_csv('submission.csv', index=False)
```

